

# MODUL PRAKTIKUM DEEP LEARNING

## Modul 8: Attention Mechanism and Transformer

Program Studi Sains Data

Fakultas Sains

Institut Teknologi Sumatera

Tahun 2025

## Daftar Isi

<b>1 Tujuan Pembelajaran</b>	<b>4</b>
<b>2 Teori dan Konsep</b>	<b>4</b>
2.1 Pengantar Attention Mechanism . . . . .	4
2.1.1 Motivasi . . . . .	4
2.2 Queries, Keys, dan Values . . . . .	4
2.2.1 Definisi Formal . . . . .	5
2.2.2 Komponen Attention . . . . .	5
2.2.3 Normalisasi Attention Weights . . . . .	5
2.3 Attention Pooling by Similarity . . . . .	5
2.3.1 Nadaraya-Watson Kernel Regression . . . . .	5
2.4 Attention Scoring Functions . . . . .	5
2.4.1 Scaled Dot-Product Attention . . . . .	6
2.4.2 Additive Attention (Bahdanau Attention) . . . . .	6
2.5 Bahdanau Attention Mechanism . . . . .	6
2.5.1 Arsitektur . . . . .	6
2.6 Multi-Head Attention . . . . .	7
2.6.1 Formulasi . . . . .	7
2.6.2 Keuntungan Multi-Head . . . . .	7
2.7 Self-Attention dan Positional Encoding . . . . .	7
2.7.1 Self-Attention . . . . .	7
2.7.2 Positional Encoding . . . . .	7
2.8 Transformer Architecture . . . . .	8
2.8.1 Arsitektur Lengkap . . . . .	8
2.8.2 Position-wise Feed-Forward Network . . . . .	8
2.8.3 Layer Normalization . . . . .	8
2.8.4 Masked Attention . . . . .	9
2.8.5 Kompleksitas Komputasi . . . . .	9
<b>3 Percobaan</b>	<b>9</b>
3.1 Percobaan 1: Implementasi Attention Scoring Functions . . . . .	9
3.1.1 Tujuan . . . . .	9
3.1.2 Langkah Percobaan . . . . .	9
3.1.3 Analisis . . . . .	12
3.2 Percobaan 2: Multi-Head Attention . . . . .	12
3.2.1 Tujuan . . . . .	12
3.2.2 Langkah Percobaan . . . . .	12
3.2.3 Analisis . . . . .	15
3.3 Percobaan 3: Positional Encoding . . . . .	15
3.3.1 Tujuan . . . . .	15
3.3.2 Langkah Percobaan . . . . .	15
3.3.3 Analisis . . . . .	17
3.4 Percobaan 4: Transformer Encoder Layer . . . . .	17
3.4.1 Tujuan . . . . .	17
3.4.2 Langkah Percobaan . . . . .	17
3.4.3 Analisis . . . . .	19
3.5 Percobaan 5: Sequence-to-Sequence dengan Attention . . . . .	19

3.5.1	Tujuan . . . . .	19
3.5.2	Dataset . . . . .	19
3.5.3	Langkah Percobaan . . . . .	20
3.5.4	Analisis . . . . .	24
<b>4</b>	<b>Latihan</b>	<b>24</b>
4.1	Latihan 1: Attention Visualization . . . . .	24
4.2	Latihan 2: Comparing Attention Mechanisms . . . . .	24
4.3	Latihan 3: Multi-Head Attention Analysis . . . . .	25
4.4	Latihan 4: Positional Encoding Variants . . . . .	25
4.5	Latihan 5: Mini Transformer untuk Text Classification . . . . .	26
4.6	Latihan 6: Masked Self-Attention . . . . .	26
<b>5</b>	<b>Dataset Rekomendasi</b>	<b>27</b>
5.1	Text Classification . . . . .	27
5.2	Machine Translation . . . . .	27
5.3	Sequence Labeling . . . . .	28
5.4	Question Answering . . . . .	28
5.5	Indonesian Language Datasets . . . . .	28
<b>6</b>	<b>Tips Implementasi</b>	<b>29</b>
6.1	Memory Optimization . . . . .	29
6.2	Training Best Practices . . . . .	29
6.3	Debugging Strategies . . . . .	29
<b>7</b>	<b>Evaluasi dan Penilaian</b>	<b>29</b>
7.1	Komponen Penilaian . . . . .	29
7.2	Kriteria Penilaian Laporan . . . . .	30
<b>8</b>	<b>Referensi</b>	<b>30</b>
<b>9</b>	<b>Lampiran</b>	<b>31</b>
9.1	Lampiran A: Cheat Sheet Formulas . . . . .	31
9.1.1	Attention Mechanisms . . . . .	31
9.1.2	Positional Encoding . . . . .	31
9.1.3	Layer Normalization . . . . .	31
9.2	Lampiran B: Common Hyperparameters . . . . .	31
9.3	Lampiran C: Troubleshooting Guide . . . . .	31

# 1 Tujuan Pembelajaran

Setelah mengikuti praktikum ini, mahasiswa diharapkan mampu:

1. Memahami konsep dasar attention mechanism dalam deep learning
2. Mengimplementasikan queries, keys, dan values dalam attention
3. Memahami dan mengimplementasikan attention pooling
4. Mengimplementasikan berbagai attention scoring functions
5. Memahami mekanisme Bahdanau attention
6. Mengimplementasikan multi-head attention
7. Memahami self-attention dan positional encoding
8. Membangun arsitektur Transformer lengkap

# 2 Teori dan Konsep

## 2.1 Pengantar Attention Mechanism

Attention mechanism adalah inovasi penting dalam deep learning yang memungkinkan model untuk fokus secara dinamis pada bagian-bagian berbeda dari input saat menghasilkan output. Berbeda dengan arsitektur RNN tradisional yang memampatkan seluruh input menjadi vektor konteks tetap, attention memungkinkan model untuk "mengakses" seluruh input sequence pada setiap langkah decoding.

### 2.1.1 Motivasi

Pada model sequence-to-sequence tradisional, encoder mengompresi seluruh input menjadi satu vektor konteks dengan dimensi tetap. Hal ini menjadi bottleneck untuk sequence yang panjang karena:

- Informasi penting bisa hilang dalam kompresi
- Model kesulitan mengingat dependensi jangka panjang
- Tidak ada mekanisme untuk fokus pada bagian relevan dari input

## 2.2 Queries, Keys, dan Values

Attention mechanism dapat dipandang sebagai operasi database yang terdiri dari tiga komponen utama:

### 2.2.1 Definisi Formal

Misalkan kita memiliki database  $\mathcal{D} = \{(k_1, v_1), \dots, (k_n, v_n)\}$  yang terdiri dari  $n$  pasangan key-value, dan sebuah query  $q$ . Attention dapat didefinisikan sebagai:

$$\text{Attention}(q, \mathcal{D}) = \sum_{i=1}^n \alpha(q, k_i) v_i \quad (1)$$

dimana  $\alpha(q, k_i)$  adalah attention weight yang menentukan seberapa besar kontribusi dari value  $v_i$ .

### 2.2.2 Komponen Attention

- **Query (q):** Representasi dari elemen yang sedang diproses, yang akan "mencari" informasi relevan
- **Keys (k):** Representasi dari semua elemen yang tersedia, digunakan untuk matching dengan query
- **Values (v):** Informasi aktual yang akan diagregasi berdasarkan attention weights

### 2.2.3 Normalisasi Attention Weights

Untuk memastikan attention weights membentuk distribusi probabilitas, kita normalisasi menggunakan softmax:

$$\alpha(q, k_i) = \frac{\exp(s(q, k_i))}{\sum_{j=1}^n \exp(s(q, k_j))} \quad (2)$$

dimana  $s(q, k_i)$  adalah scoring function yang mengukur kompatibilitas antara query dan key.

## 2.3 Attention Pooling by Similarity

Attention pooling mengagregasi values berdasarkan kesamaan antara queries dan keys.

### 2.3.1 Nadaraya-Watson Kernel Regression

Sebagai contoh klasik, Nadaraya-Watson estimator menggunakan kernel similarity:

$$f(q) = \sum_{i=1}^n \frac{K(q, k_i)}{\sum_{j=1}^n K(q, k_j)} v_i \quad (3)$$

**Kernel Functions yang Umum:**

$$K_{\text{Gaussian}}(q, k) = \exp\left(-\frac{1}{2}\|q - k\|^2\right) \quad (4)$$

$$K_{\text{Boxcar}}(q, k) = \mathbb{1}(\|q - k\| \leq 1) \quad (5)$$

$$K_{\text{Epanechnikov}}(q, k) = \max(0, 1 - \|q - k\|) \quad (6)$$

## 2.4 Attention Scoring Functions

Scoring function menentukan seberapa baik query dan key "match" satu sama lain.

### 2.4.1 Scaled Dot-Product Attention

Ini adalah scoring function paling populer dalam Transformer:

$$s(q, k_i) = \frac{q^T k_i}{\sqrt{d}} \quad (7)$$

dimana  $d$  adalah dimensi dari query dan key. Pembagian dengan  $\sqrt{d}$  mencegah nilai dot product menjadi terlalu besar.

**Attention Lengkap:**

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d}} \right) V \quad (8)$$

dimana:

- $Q \in \mathbb{R}^{n \times d}$ : matrix dari  $n$  queries
- $K \in \mathbb{R}^{m \times d}$ : matrix dari  $m$  keys
- $V \in \mathbb{R}^{m \times d_v}$ : matrix dari  $m$  values

### 2.4.2 Additive Attention (Bahdanau Attention)

Alternatif untuk dot-product, menggunakan feed-forward network:

$$s(q, k) = w_v^T \tanh(W_q q + W_k k) \quad (9)$$

dimana  $W_q \in \mathbb{R}^{h \times d_q}$ ,  $W_k \in \mathbb{R}^{h \times d_k}$ , dan  $w_v \in \mathbb{R}^h$  adalah parameter yang dapat dipelajari.

## 2.5 Bahdanau Attention Mechanism

Bahdanau attention adalah salah satu implementasi attention pertama untuk sequence-to-sequence learning.

### 2.5.1 Arsitektur

Dalam konteks machine translation:

1. Encoder menghasilkan hidden states  $h_1, \dots, h_T$  untuk input sequence
2. Pada setiap decoding step  $t'$ , decoder menggunakan:
  - Previous decoder state  $s_{t'-1}$  sebagai query
  - Semua encoder hidden states  $h_t$  sebagai keys dan values

3. Context vector dihitung:

$$c_{t'} = \sum_{t=1}^T \alpha_{t',t} h_t \quad (10)$$

4. Attention weights:

$$\alpha_{t',t} = \frac{\exp(e_{t',t})}{\sum_{j=1}^T \exp(e_{t',j})} \quad (11)$$

dimana  $e_{t',t} = s(s_{t'-1}, h_t)$  adalah alignment score

## 2.6 Multi-Head Attention

Multi-head attention memungkinkan model untuk fokus pada informasi dari berbagai representation subspaces secara paralel.

### 2.6.1 Formulasi

$$\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O \quad (12)$$

dimana setiap head dihitung sebagai:

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \quad (13)$$

**Parameter Matrices:**

- $W_i^Q \in \mathbb{R}^{d \times d_k}$ : projection matrix untuk queries
- $W_i^K \in \mathbb{R}^{d \times d_k}$ : projection matrix untuk keys
- $W_i^V \in \mathbb{R}^{d \times d_v}$ : projection matrix untuk values
- $W^O \in \mathbb{R}^{hd_v \times d}$ : output projection matrix

### 2.6.2 Keuntungan Multi-Head

1. Memungkinkan model attend ke berbagai posisi secara paralel
2. Setiap head dapat mempelajari aspek berbeda dari input
3. Meningkatkan kapasitas model tanpa menambah kompleksitas komputasi secara signifikan

## 2.7 Self-Attention dan Positional Encoding

### 2.7.1 Self-Attention

Dalam self-attention, queries, keys, dan values semuanya berasal dari sequence yang sama:

$$\text{SelfAttention}(X) = \text{Attention}(XW^Q, XW^K, XW^V) \quad (14)$$

dimana  $X \in \mathbb{R}^{n \times d}$  adalah input sequence.

### 2.7.2 Positional Encoding

Karena self-attention tidak memiliki notion of order, kita perlu menambahkan informasi posisi:

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d}}\right) \quad (15)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d}}\right) \quad (16)$$

dimana:

- $pos$  adalah posisi token dalam sequence

- $i$  adalah dimensi
- $d$  adalah dimensi model

### Keuntungan Positional Encoding:

- Dapat menangani sequence dengan panjang arbitrary
- Mengandung informasi posisi absolut dan relatif
- Dapat direpresentasikan sebagai linear projection

## 2.8 Transformer Architecture

### 2.8.1 Arsitektur Lengkap

Transformer terdiri dari encoder dan decoder stack, masing-masing dengan komponen:

#### Encoder Layer:

1. Multi-head self-attention
2. Add & Norm (residual connection + layer normalization)
3. Position-wise feed-forward network
4. Add & Norm

#### Decoder Layer:

1. Masked multi-head self-attention
2. Add & Norm
3. Multi-head cross-attention (attend to encoder output)
4. Add & Norm
5. Position-wise feed-forward network
6. Add & Norm

### 2.8.2 Position-wise Feed-Forward Network

$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2 \quad (17)$$

atau dengan aktivasi GELU:

$$\text{FFN}(x) = \text{GELU}(xW_1 + b_1)W_2 + b_2 \quad (18)$$

### 2.8.3 Layer Normalization

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} + \beta \quad (19)$$

dimana  $\mu$  dan  $\sigma^2$  adalah mean dan variance dihitung across feature dimension.

### 2.8.4 Masked Attention

Dalam decoder, kita perlu mencegah attending ke future tokens:

$$\text{mask}_{ij} = \begin{cases} 0 & \text{if } i \geq j \\ -\infty & \text{if } i < j \end{cases} \quad (20)$$

Attention dengan mask:

$$\text{Attention}_{\text{masked}}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d}} + M \right) V \quad (21)$$

### 2.8.5 Kompleksitas Komputasi

Layer Type	Complexity	Sequential	Max Path
Self-Attention	$O(n^2 \cdot d)$	$O(1)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$	$O(\log_k(n))$

Tabel 1: Perbandingan kompleksitas berbagai layer types

dimana  $n$  adalah sequence length,  $d$  adalah dimensi representasi, dan  $k$  adalah kernel size.

## 3 Percobaan

### 3.1 Percobaan 1: Implementasi Attention Scoring Functions

#### 3.1.1 Tujuan

Mengimplementasikan dan membandingkan berbagai attention scoring functions.

#### 3.1.2 Langkah Percobaan

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 def softmax(x, axis=-1):
5     """Compute softmax values"""
6     exp_x = np.exp(x - np.max(x, axis=axis, keepdims=True))
7     return exp_x / np.sum(exp_x, axis=axis, keepdims=True)
8
9 def scaled_dot_product_attention(Q, K, V, mask=None):
10    """
11    Scaled Dot-Product Attention
12
13    Args:
14        Q: Queries (batch_size, n_queries, d_k)
15        K: Keys (batch_size, n_keys, d_k)
16        V: Values (batch_size, n_keys, d_v)

```

```

17     mask: Optional mask (batch_size, n_queries, n_keys)
18
19     Returns:
20         output: Attention output
21         attention_weights: Attention weights
22     """
23
24     d_k = Q.shape[-1]
25
26     # Compute attention scores
27     scores = np.matmul(Q, K.transpose(0, 2, 1)) / np.sqrt(d_k)
28
29     # Apply mask if provided
30     if mask is not None:
31         scores = scores + (mask * -1e9)
32
33     # Compute attention weights
34     attention_weights = softmax(scores, axis=-1)
35
36     # Compute output
37     output = np.matmul(attention_weights, V)
38
39     return output, attention_weights
40
41 def additive_attention(Q, K, V, W_q, W_k, w_v):
42     """
43     Additive Attention (Bahdanau)
44
45     Args:
46         Q: Queries (batch_size, n_queries, d_q)
47         K: Keys (batch_size, n_keys, d_k)
48         V: Values (batch_size, n_keys, d_v)
49         W_q: Query projection (d_q, h)
50         W_k: Key projection (d_k, h)
51         w_v: Score projection (h, )
52
53     Returns:
54         output: Attention output
55         attention_weights: Attention weights
56     """
57
58     batch_size, n_queries, _ = Q.shape
59     n_keys = K.shape[1]
60
61     # Project queries and keys
62     Q_proj = np.matmul(Q, W_q) # (batch, n_queries, h)
63     K_proj = np.matmul(K, W_k) # (batch, n_keys, h)
64
65     # Expand dimensions for broadcasting
66     Q_exp = Q_proj[:, :, np.newaxis, :] # (batch, n_queries, 1, h)
67     K_exp = K_proj[:, np.newaxis, :, :] # (batch, 1, n_keys, h)
68
69     # Compute scores

```

```

68     features = np.tanh(Q_exp + K_exp) # (batch, n_queries, n_keys,
69     h)
70     scores = np.dot(features, w_v)      # (batch, n_queries, n_keys)
71
72     # Compute attention weights
73     attention_weights = softmax(scores, axis=-1)
74
75     # Compute output
76     output = np.matmul(attention_weights, V)
77
78     return output, attention_weights
79
80 # Test implementations
81 np.random.seed(42)
82
83 # Generate sample data
84 batch_size, n_queries, n_keys = 2, 4, 6
85 d_k, d_v, h = 8, 10, 16
86
87 Q = np.random.randn(batch_size, n_queries, d_k)
88 K = np.random.randn(batch_size, n_keys, d_k)
89 V = np.random.randn(batch_size, n_keys, d_v)
90
91 # Test Scaled Dot-Product Attention
92 output_sdp, weights_sdp = scaled_dot_product_attention(Q, K, V)
93 print("Scaled Dot-Product Attention:")
94 print(f"Output shape: {output_sdp.shape}")
95 print(f"Attention weights shape: {weights_sdp.shape}")
96 print(f"Attention weights sum: {weights_sdp.sum(axis=-1)}")
97
98 # Test Additive Attention
99 W_q = np.random.randn(d_k, h) * 0.1
100 W_k = np.random.randn(d_k, h) * 0.1
101 w_v = np.random.randn(h) * 0.1
102
103 output_add, weights_add = additive_attention(Q, K, V, W_q, W_k, w_v)
104 print("\nAdditive Attention:")
105 print(f"Output shape: {output_add.shape}")
106 print(f"Attention weights shape: {weights_add.shape}")
107 print(f"Attention weights sum: {weights_add.sum(axis=-1)}")
108
109 # Visualize attention weights
110 fig, axes = plt.subplots(1, 2, figsize=(12, 4))
111 im1 = axes[0].imshow(weights_sdp[0], cmap='Blues', aspect='auto')
112 axes[0].set_title('Scaled Dot-Product Attention Weights')
113 axes[0].set_xlabel('Keys')
114 axes[0].set_ylabel('Queries')
115 plt.colorbar(im1, ax=axes[0])
116

```

```

117 im2 = axes[1].imshow(weights_add[0], cmap='Blues', aspect='auto')
118 axes[1].set_title('Additive Attention Weights')
119 axes[1].set_xlabel('Keys')
120 axes[1].set_ylabel('Queries')
121 plt.colorbar(im2, ax=axes[1])
122
123 plt.tight_layout()
124 plt.savefig('attention_comparison.png', dpi=300, bbox_inches='tight')
125 plt.show()

```

Listing 1: Implementasi Scoring Functions

### 3.1.3 Analisis

1. Bandingkan pattern attention weights dari kedua metode
2. Perhatikan bahwa sum dari attention weights untuk setiap query adalah 1
3. Scaled dot-product lebih efisien secara komputasi
4. Additive attention lebih fleksibel untuk dimensi yang berbeda

## 3.2 Percobaan 2: Multi-Head Attention

### 3.2.1 Tujuan

Mengimplementasikan multi-head attention mechanism.

### 3.2.2 Langkah Percobaan

```

1 class MultiHeadAttention:
2     def __init__(self, d_model, num_heads):
3         """
4             Multi-Head Attention
5
6         Args:
7             d_model: Dimension of the model
8             num_heads: Number of attention heads
9         """
10        assert d_model % num_heads == 0, "d_model must be divisible
11        by num_heads"
12
13        self.d_model = d_model
14        self.num_heads = num_heads
15        self.d_k = d_model // num_heads
16
17        # Initialize projection matrices
18        self.W_q = np.random.randn(d_model, d_model) * 0.1
19        self.W_k = np.random.randn(d_model, d_model) * 0.1
20        self.W_v = np.random.randn(d_model, d_model) * 0.1
21        self.W_o = np.random.randn(d_model, d_model) * 0.1

```

```

21
22     def split_heads(self, x):
23         """
24             Split the last dimension into (num_heads, d_k)
25
26         Args:
27             x: Input tensor (batch_size, seq_len, d_model)
28
29         Returns:
30             Reshaped tensor (batch_size, num_heads, seq_len, d_k)
31         """
32         batch_size, seq_len, _ = x.shape
33         x = x.reshape(batch_size, seq_len, self.num_heads, self.d_k)
34             )
35         return x.transpose(0, 2, 1, 3)
36
37     def combine_heads(self, x):
38         """
39             Combine heads back to original shape
40
41         Args:
42             x: Input tensor (batch_size, num_heads, seq_len, d_k)
43
44         Returns:
45             Combined tensor (batch_size, seq_len, d_model)
46         """
47         batch_size, _, seq_len, _ = x.shape
48         x = x.transpose(0, 2, 1, 3)
49         return x.reshape(batch_size, seq_len, self.d_model)
50
51     def forward(self, Q, K, V, mask=None):
52         """
53             Forward pass of multi-head attention
54
55         Args:
56             Q: Queries (batch_size, seq_len_q, d_model)
57             K: Keys (batch_size, seq_len_k, d_model)
58             V: Values (batch_size, seq_len_v, d_model)
59             mask: Optional mask
60
61         Returns:
62             output: Attention output
63             attention_weights: Attention weights for all heads
64         """
65         batch_size = Q.shape[0]
66
67         # Linear projections
68         Q = np.matmul(Q, self.W_q)
69         K = np.matmul(K, self.W_k)
70         V = np.matmul(V, self.W_v)

```

```

71     # Split into multiple heads
72     Q = self.split_heads(Q)  # (batch, num_heads, seq_len_q,
73                           # d_k)
73     K = self.split_heads(K)  # (batch, num_heads, seq_len_k,
74                           # d_k)
74     V = self.split_heads(V)  # (batch, num_heads, seq_len_v,
75                           # d_k)

75
76     # Scaled dot-product attention
77     scores = np.matmul(Q, K.transpose(0, 1, 3, 2)) / np.sqrt(
78                           self.d_k)

78
79     if mask is not None:
80         scores = scores + (mask * -1e9)

81
82     attention_weights = softmax(scores, axis=-1)
83     attention_output = np.matmul(attention_weights, V)

84
85     # Combine heads
86     output = self.combine_heads(attention_output)

87
88     # Final linear projection
89     output = np.matmul(output, self.W_o)

90
91     return output, attention_weights

92

93 # Test Multi-Head Attention
94 np.random.seed(42)

95
96 d_model, num_heads = 512, 8
97 batch_size, seq_len = 2, 10

98
99 mha = MultiHeadAttention(d_model, num_heads)

100
101 Q = np.random.randn(batch_size, seq_len, d_model)
102 K = np.random.randn(batch_size, seq_len, d_model)
103 V = np.random.randn(batch_size, seq_len, d_model)

104
105 output, attention_weights = mha.forward(Q, K, V)

106
107 print(f"Input shape: {Q.shape}")
108 print(f"Output shape: {output.shape}")
109 print(f"Attention weights shape: {attention_weights.shape}")
110 print(f"Number of heads: {num_heads}")
111 print(f"d_k per head: {mha.d_k}")

112
113 # Visualize attention patterns for different heads
114 fig, axes = plt.subplots(2, 4, figsize=(16, 8))
115 axes = axes.flatten()

116
117 for i in range(num_heads):

```

```
118     im = axes[i].imshow(attention_weights[0, i], cmap='viridis',
119                          aspect='auto')
120     axes[i].set_title(f'Head {i+1}')
121     axes[i].set_xlabel('Keys')
122     axes[i].set_ylabel('Queries')
123     plt.colorbar(im, ax=axes[i])
124
125 plt.tight_layout()
126 plt.savefig('multihead_attention.png', dpi=300, bbox_inches='tight')
127 plt.show()
```

Listing 2: Multi-Head Attention Implementation

### 3.2.3 Analisis

1. Perhatikan bagaimana setiap head memiliki pattern attention yang berbeda
  2. Setiap head dapat fokus pada aspek berbeda dari input
  3. Multi-head attention meningkatkan ekspresivitas model
  4. Total parameter meningkat tetapi kompleksitas per head berkurang

### 3.3 Percobaan 3: Positional Encoding

### 3.3.1 Tujuan

Mengimplementasikan dan memvisualisasikan positional encoding.

### 3.3.2 Langkah Percobaan

```
1 def positional_encoding(max_len, d_model):
2     """
3         Generate positional encoding
4
5     Args:
6         max_len: Maximum sequence length
7         d_model: Dimension of the model
8
9     Returns:
10        pos_encoding: Positional encoding matrix (max_len, d_model)
11    """
12    pos_encoding = np.zeros((max_len, d_model))
13    position = np.arange(0, max_len)[:, np.newaxis]
14    div_term = np.exp(np.arange(0, d_model, 2) *
15                      -(np.log(10000.0) / d_model))
16
17    # Apply sin to even indices
18    pos_encoding[:, 0::2] = np.sin(position * div_term)
19
20    # Apply cos to odd indices
21    pos_encoding[:, 1::2] = np.cos(position * div_term)
22
23    return pos_encoding
```

```
21 pos_encoding[:, 1::2] = np.cos(position * div_term)
22
23 return pos_encoding
24
25 # Generate positional encoding
26 max_len, d_model = 100, 512
27 pos_enc = positional_encoding(max_len, d_model)
28
29 print(f"Positional_encoding shape: {pos_enc.shape}")
30
31 # Visualize positional encoding
32 plt.figure(figsize=(15, 5))
33
34 # Plot full positional encoding
35 plt.subplot(1, 2, 1)
36 plt.imshow(pos_enc.T, cmap='RdBu', aspect='auto')
37 plt.colorbar()
38 plt.xlabel('Position')
39 plt.ylabel('Dimension')
40 plt.title('Positional_Encoding_Heatmap')
41
42 # Plot specific dimensions over positions
43 plt.subplot(1, 2, 2)
44 dims_to_plot = [0, 1, 64, 65, 128, 129]
45 for dim in dims_to_plot:
46     plt.plot(pos_enc[:, dim], label=f'Dim_{dim}')
47 plt.xlabel('Position')
48 plt.ylabel('Value')
49 plt.title('Positional_Encoding_Values_for_Different_Dimensions')
50 plt.legend()
51 plt.grid(True, alpha=0.3)
52
53 plt.tight_layout()
54 plt.savefig('positional_encoding.png', dpi=300, bbox_inches='tight')
55 plt.show()
56
57 # Test relative position property
58 def demonstrate_relative_position(pos_enc, offset=5):
59     """
60         Demonstrate that relative positions can be represented
61         as linear transformations
62     """
63     position_i = 10
64     position_j = position_i + offset
65
66     # Get encodings
67     pe_i = pos_enc[position_i]
68     pe_j = pos_enc[position_j]
69
70     # Compute similarity
```

```

71     similarity = np.dot(pe_i, pe_j) / (np.linalg.norm(pe_i) *
72                                         np.linalg.norm(pe_j))
73
74     print(f"\nRelative Position Analysis:")
75     print(f"Position {position_i} to {position_j} (offset={offset})")
76     print(f"Cosine similarity: {similarity:.4f}")
77
78     return similarity
79
80 # Test different offsets
81 offsets = range(0, 20)
82 similarities = [demonstrate_relative_position(pos_enc, offset)
83                  for offset in offsets]
84
85 plt.figure(figsize=(10, 5))
86 plt.plot(offsets, similarities, 'o-')
87 plt.xlabel('Position Offset')
88 plt.ylabel('Cosine Similarity')
89 plt.title('Positional Encoding Similarity vs Position Offset')
90 plt.grid(True, alpha=0.3)
91 plt.savefig('pe_similarity.png', dpi=300, bbox_inches='tight')
92 plt.show()

```

Listing 3: Positional Encoding

### 3.3.3 Analisis

1. Positional encoding menambahkan informasi posisi ke embeddings
2. Dimensi dengan frekuensi rendah berubah lambat across positions
3. Dimensi dengan frekuensi tinggi berubah cepat
4. Model dapat mempelajari relative positions dari pattern ini

## 3.4 Percobaan 4: Transformer Encoder Layer

### 3.4.1 Tujuan

Mengimplementasikan complete Transformer encoder layer.

### 3.4.2 Langkah Percobaan

```

1 class LayerNormalization:
2     def __init__(self, features, eps=1e-6):
3         self.eps = eps
4         self.gamma = np.ones(features)
5         self.beta = np.zeros(features)
6
7     def forward(self, x):
8         mean = np.mean(x, axis=-1, keepdims=True)

```

```

9      std = np.std(x, axis=-1, keepdims=True)
10     return self.gamma * (x - mean) / (std + self.eps) + self.
11         beta
12
13 class PositionwiseFeedForward:
14     def __init__(self, d_model, d_ff):
15         self.W1 = np.random.randn(d_model, d_ff) * 0.1
16         self.b1 = np.zeros(d_ff)
17         self.W2 = np.random.randn(d_ff, d_model) * 0.1
18         self.b2 = np.zeros(d_model)
19
20     def forward(self, x):
21         # First layer with ReLU
22         hidden = np.maximum(0, np.matmul(x, self.W1) + self.b1)
23         # Second layer
24         output = np.matmul(hidden, self.W2) + self.b2
25         return output
26
27 class TransformerEncoderLayer:
28     def __init__(self, d_model, num_heads, d_ff, dropout=0.1):
29         self.mha = MultiHeadAttention(d_model, num_heads)
30         self.ffn = PositionwiseFeedForward(d_model, d_ff)
31         self.layernorm1 = LayerNormalization(d_model)
32         self.layernorm2 = LayerNormalization(d_model)
33         self.dropout = dropout
34
35     def forward(self, x, mask=None):
36         # Multi-head attention with residual connection and layer
37             norm
38         attn_output, attn_weights = self.mha.forward(x, x, x, mask)
39
40         # Dropout (simulated by scaling)
41         attn_output = attn_output * (1 - self.dropout)
42
43         # Add & Norm
44         x = self.layernorm1.forward(x + attn_output)
45
46         # Feed-forward with residual connection and layer norm
47         ffn_output = self.ffn.forward(x)
48
49         # Dropout
50         ffn_output = ffn_output * (1 - self.dropout)
51
52         # Add & Norm
53         output = self.layernorm2.forward(x + ffn_output)
54
55     # Test Transformer Encoder Layer
56     np.random.seed(42)
57

```

```

58 d_model, num_heads, d_ff = 512, 8, 2048
59 batch_size, seq_len = 2, 10
60
61 encoder_layer = TransformerEncoderLayer(d_model, num_heads, d_ff)
62
63 # Generate input with positional encoding
64 x = np.random.randn(batch_size, seq_len, d_model)
65 pos_enc = positional_encoding(seq_len, d_model)
66 x = x + pos_enc
67
68 # Forward pass
69 output, attn_weights = encoder_layer.forward(x)
70
71 print(f"Input shape: {x.shape}")
72 print(f"Output shape: {output.shape}")
73 print(f"Attention weights shape: {attn_weights.shape}")
74
75 # Visualize attention for first sample
76 plt.figure(figsize=(10, 8))
77 avg_attn = np.mean(attn_weights[0], axis=0)
78 plt.imshow(avg_attn, cmap='viridis', aspect='auto')
79 plt.colorbar(label='Attention Weight')
80 plt.xlabel('Key Position')
81 plt.ylabel('Query Position')
82 plt.title('Average Attention Weights Across All Heads')
83 plt.savefig('encoder_attention.png', dpi=300, bbox_inches='tight')
84 plt.show()

```

Listing 4: Transformer Encoder Layer

### 3.4.3 Analisis

1. Encoder layer menggabungkan multi-head attention dan feed-forward network
2. Residual connections membantu gradient flow
3. Layer normalization menstabilkan training
4. Setiap token dapat attend ke semua token lainnya

## 3.5 Percobaan 5: Sequence-to-Sequence dengan Attention

### 3.5.1 Tujuan

Mengimplementasikan complete sequence-to-sequence model dengan attention untuk machine translation.

### 3.5.2 Dataset

Gunakan dataset English-French translation pairs dari Tatoeba Project atau Manythings.org.

### 3.5.3 Langkah Percobaan

```

1 import pandas as pd
2 from collections import Counter
3 import re
4
5 class Vocabulary:
6     def __init__(self, freq_threshold=2):
7         self.itos = {0: "<PAD>", 1: "<SOS>", 2: "<EOS>", 3: "<UNK>"}
8         self.stoi = {"<PAD>": 0, "<SOS>": 1, "<EOS>": 2, "<UNK>": 3}
9         self.freq_threshold = freq_threshold
10
11    def build_vocabulary(self, sentence_list):
12        frequencies = Counter()
13        idx = 4
14
15        for sentence in sentence_list:
16            for word in sentence.split():
17                frequencies[word] += 1
18
19                if frequencies[word] == self.freq_threshold:
20                    self.stoi[word] = idx
21                    self.itos[idx] = word
22                    idx += 1
23
24    def numericalize(self, text):
25        tokenized = text.split()
26        return [self.stoi.get(token, self.stoi["<UNK>"]))
27                    for token in tokenized]
28
29    def __len__(self):
30        return len(self.itos)
31
32    def preprocess_text(text):
33        """Simple text preprocessing"""
34        text = text.lower()
35        text = re.sub(r"([?.! ,])", r"\1 ", text)
36        text = re.sub(r'["_"]+', " ", text)
37        text = text.strip()
38        return text
39
40    # Load and prepare data
41    def load_translation_data(file_path, num_examples=10000):
42        """
43            Load translation data
44            Expected format: English\tFrench
45        """
46        data = pd.read_csv(file_path, sep='\t', header=None,

```

```

47             names=[ 'english' , 'french' ] , nrows=
48             num_examples)
49
50     data[ 'english' ] = data[ 'english' ].apply(preprocess_text)
51     data[ 'french' ] = data[ 'french' ].apply(preprocess_text)
52
53     return data
54
55 # Simple Seq2Seq with Attention (using simplified encoder-decoder)
56 class SimpleSeq2SeqAttention:
57     def __init__(self, src_vocab_size, tgt_vocab_size,
58                  embed_size, hidden_size, num_heads=4):
59         self.embed_size = embed_size
60         self.hidden_size = hidden_size
61
62         # Embeddings
63         self.src_embedding = np.random.randn(src_vocab_size,
64                                              embed_size) * 0.1
64         self.tgt_embedding = np.random.randn(tgt_vocab_size,
65                                              embed_size) * 0.1
65
66
67         # Encoder (simplified as linear transformation)
68         self.encoder_transform = np.random.randn(embed_size,
69                                              hidden_size) * 0.1
69
70
71         # Attention mechanism
72         self.attention = MultiHeadAttention(hidden_size, num_heads)
73
74
75         # Decoder
76         self.decoder_transform = np.random.randn(hidden_size,
77                                              hidden_size) * 0.1
77         self.output_projection = np.random.randn(hidden_size,
78                                              tgt_vocab_size) *
78                                              0.1
79
80     def encode(self, src_seq):
81         """
82             Encode source sequence
83
84             Args:
85                 src_seq: Source sequence indices (batch_size, src_len)
86
87             Returns:
88                 Encoder outputs (batch_size, src_len, hidden_size)
89             """
90
91         # Get embeddings
92         src_embed = self.src_embedding[src_seq]
93
94         # Add positional encoding
95         pos_enc = positional_encoding(src_seq.shape[1], self.
95                                         embed_size)

```

```

95     src_embed = src_embed + pos_enc
96
97     # Transform to hidden size
98     encoder_output = np.matmul(src_embed, self.
99         encoder_transform)
100
101    return encoder_output
102
103    def decode_step(self, tgt_input, encoder_output, mask=None):
104        """
105            Single decoding step
106
107            Args:
108                tgt_input: Target input (batch_size, tgt_len,
109                    hidden_size)
110                encoder_output: Encoder outputs (batch_size, src_len,
111                    hidden_size)
112                mask: Optional attention mask
113
114            Returns:
115                decoder_output: Predictions (batch_size, tgt_len,
116                    tgt_vocab_size)
117                attention_weights: Attention weights
118        """
119
120        # Cross-attention
121        attended, attention_weights = self.attention.forward(
122            tgt_input, encoder_output, encoder_output, mask
123        )
124
125        # Decoder transformation
126        decoder_hidden = np.matmul(attended, self.decoder_transform
127            )
128        decoder_hidden = np.maximum(0, decoder_hidden) # ReLU
129
130        # Output projection
131        decoder_output = np.matmul(decoder_hidden, self.
132            output_projection)
133
134        return decoder_output, attention_weights
135
136    def predict(self, src_seq, max_len=20):
137        """
138            Generate translation
139
140            Args:
141                src_seq: Source sequence (batch_size, src_len)
142                max_len: Maximum target length
143
144            Returns:
145                predictions: Generated sequence
146        """

```

```

140     # Encode
141     encoder_output = self.encode(src_seq)
142
143     batch_size = src_seq.shape[0]
144     predictions = np.ones((batch_size, 1), dtype=int) # Start
145     with <SOS>
146
146     for _ in range(max_len):
147         # Get target embeddings
148         tgt_embed = self.tgt_embedding[predictions]
149         pos_enc = positional_encoding(predictions.shape[1],
150                                       self.embed_size)
151         tgt_embed = tgt_embed + pos_enc[:predictions.shape[1]]
152
153         # Transform to hidden size
154         tgt_hidden = np.matmul(tgt_embed, self.
155                               encoder_transform)
156
156         # Decode
157         output, _ = self.decode_step(tgt_hidden, encoder_output
158                                       )
159
159         # Get next token
160         next_token = np.argmax(output[:, -1, :], axis=-1,
161                               keepdims=True)
161         predictions = np.concatenate([predictions, next_token],
162                                       axis=1)
162
163         # Stop if all sequences generated <EOS>
164         if np.all(next_token == 2): # 2 is <EOS>
165             break
166
167     return predictions
168
169 # Example usage (with dummy data for demonstration)
170 print("Creating vocabularies and model...")
171
172 # Create dummy data
173 src_vocab_size, tgt_vocab_size = 5000, 5000
174 embed_size, hidden_size = 256, 512
175
176 model = SimpleSeq2SeqAttention(src_vocab_size, tgt_vocab_size,
177                                 embed_size, hidden_size)
178
179 # Test with random input
180 batch_size, src_len = 2, 10
181 src_seq = np.random.randint(0, src_vocab_size, (batch_size, src_len
182                               ))
183 print(f"Source sequence shape: {src_seq.shape}")
184

```

```

185 # Encode
186 encoder_output = model.encode(src_seq)
187 print(f"Encoder output shape: {encoder_output.shape}")
188
189 # Generate translation
190 predictions = model.predict(src_seq, max_len=15)
191 print(f"Predictions shape: {predictions.shape}")
192 print(f"Sample prediction: {predictions[0]}")

```

Listing 5: Seq2Seq dengan Attention

### 3.5.4 Analisis

1. Model dapat menangani sequences dengan panjang berbeda
2. Attention memungkinkan decoder fokus pada bagian relevan dari input
3. Positional encoding penting untuk mempertahankan informasi urutan
4. Model dapat di-extend dengan multiple layers untuk performa lebih baik

## 4 Latihan

### 4.1 Latihan 1: Attention Visualization

**Tugas:**

1. Implementasikan fungsi untuk memvisualisasikan attention weights sebagai heatmap
2. Gunakan dataset teks sederhana (misalnya kalimat dalam bahasa Indonesia)
3. Visualisasikan bagaimana attention weights berubah untuk berbagai query positions
4. Analisis pattern yang muncul

**Pertanyaan:**

1. Apa perbedaan pattern attention antara short-range dan long-range dependencies?
2. Bagaimana attention weights berubah untuk kata-kata yang berbeda jenis (noun, verb, dll)?

### 4.2 Latihan 2: Comparing Attention Mechanisms

**Tugas:**

1. Implementasikan tiga attention mechanisms:
  - Scaled Dot-Product Attention
  - Additive Attention
  - Multiplicative Attention (variant dari dot-product)

2. Bandingkan kompleksitas komputasi dan memory usage
3. Test pada sequence dengan panjang berbeda (10, 50, 100, 500)
4. Plot hasil comparison

**Deliverables:**

- Code implementation
- Computational complexity analysis
- Performance comparison plots
- Written analysis (min. 500 words)

### 4.3 Latihan 3: Multi-Head Attention Analysis

**Tugas:**

1. Implementasikan multi-head attention dengan jumlah heads yang berbeda (1, 2, 4, 8, 16)
2. Analisis bagaimana setiap head mempelajari pattern yang berbeda
3. Visualisasikan attention patterns untuk setiap head
4. Compute attention entropy untuk mengukur "fokus" dari attention

**Attention Entropy Formula:**

$$H = - \sum_{i=1}^n \alpha_i \log(\alpha_i) \quad (22)$$

**Pertanyaan Analisis:**

1. Apakah semua heads mempelajari pattern yang unik?
2. Bagaimana jumlah heads mempengaruhi kapasitas model?
3. Trade-off apa yang ada antara jumlah heads dan performa?

### 4.4 Latihan 4: Positional Encoding Variants

**Tugas:**

1. Implementasikan tiga jenis positional encoding:
  - Sinusoidal (original Transformer)
  - Learned positional embeddings
  - Relative positional encoding
2. Bandingkan performa pada task sequence classification
3. Analisis generalisasi ke sequence lengths yang tidak terlihat saat training

**Relative Positional Encoding:**

$$\text{RelPE}(i, j) = \text{clip}(i - j, -k, k) \quad (23)$$

dimana  $k$  adalah maximum relative distance.

## 4.5 Latihan 5: Mini Transformer untuk Text Classification

**Tugas:** Implementasikan complete Transformer encoder untuk sentiment analysis.

**Spesifikasi:**

- Dataset: IMDb Movie Reviews atau AG News
- Architecture:
  - Embedding layer ( $d_{\text{model}} = 256$ )
  - Positional encoding
  - 2 Transformer encoder layers
  - 4 attention heads
  - Feed-forward dimension = 1024
  - Classification head
- Training: 10 epochs
- Evaluation metrics: Accuracy, F1-score

**Deliverables:**

1. Complete code implementation
2. Training curves (loss and accuracy)
3. Attention visualization for sample predictions
4. Error analysis
5. Comparison dengan baseline model (e.g., LSTM)

## 4.6 Latihan 6: Masked Self-Attention

**Tugas:**

1. Implementasikan masked self-attention untuk language modeling
2. Generate causal attention mask
3. Implementasikan autoregressive generation
4. Test pada dataset teks sederhana

**Causal Mask:**

$$M_{ij} = \begin{cases} 0 & \text{if } i \geq j \\ -\infty & \text{if } i < j \end{cases} \quad (24)$$

**Tasks:**

1. Implement causal mask generation
2. Verify that future tokens are not attended
3. Generate text sequences autoregressively
4. Measure generation quality (perplexity)

## 5 Dataset Rekomendasi

### 5.1 Text Classification

#### 1. IMDb Movie Reviews

- URL: <https://ai.stanford.edu/~amaas/data/sentiment/>
- Size: 50,000 reviews
- Task: Binary sentiment classification
- Format: Text files

#### 2. AG News

- URL: <https://www.kaggle.com/datasets/amananandrai/ag-news-classification-dataset>
- Size: 120,000 training, 7,600 test
- Task: 4-class news categorization
- Categories: World, Sports, Business, Sci/Tech

#### 3. 20 Newsgroups

- URL: <http://qwone.com/~jason/20Newsgroups/>
- Size: 20,000 documents
- Task: Multi-class classification (20 classes)
- Format: Text documents

### 5.2 Machine Translation

#### 1. Tatoeba Translation Pairs

- URL: <https://tatoeba.org/en/downloads>
- Languages: 300+ languages
- Size: Varies by language pair
- Format: Tab-separated values

#### 2. WMT Translation Datasets

- URL: <https://www.statmt.org/wmt22/>
- Task: Various language pairs
- Size: Millions of sentence pairs
- Quality: Professional translations

#### 3. Multi30k

- URL: <https://github.com/multi30k/dataset>
- Languages: English, German, French, Czech
- Size: 31,000 image descriptions
- Task: Multilingual image captioning

### 5.3 Sequence Labeling

#### 1. CoNLL 2003 NER

- URL: <https://www.clips.uantwerpen.be/conll2003/ner/>
- Task: Named Entity Recognition
- Tags: PER, LOC, ORG, MISC
- Format: CoNLL format

#### 2. Universal Dependencies

- URL: <https://universaldependencies.org/>
- Languages: 100+ languages
- Task: POS tagging, dependency parsing
- Format: CoNLL-U

### 5.4 Question Answering

#### 1. SQuAD 2.0

- URL: <https://rajpurkar.github.io/SQuAD-explorer/>
- Size: 100,000+ questions
- Task: Extractive QA
- Format: JSON

#### 2. Natural Questions

- URL: <https://ai.google.com/research/NaturalQuestions>
- Size: 300,000+ questions
- Source: Real Google searches
- Format: JSON

### 5.5 Indonesian Language Datasets

#### 1. IndoNLU

- URL: <https://github.com/indobenchmark/indonlu>
- Tasks: Sentiment, NER, QA, etc.
- Language: Indonesian
- Size: Varies by task

#### 2. Indonesian News Dataset

- URL: <https://www.kaggle.com/datasets/scolianni/indonesian-news-dataset>
- Size: 100,000+ articles
- Task: Classification, summarization
- Categories: Multiple news categories

## 6 Tips Implementasi

### 6.1 Memory Optimization

1. **Gradient Checkpointing:** Trade compute for memory
2. **Mixed Precision Training:** Use float16 untuk forward pass
3. **Gradient Accumulation:** Simulate larger batch sizes
4. **Efficient Attention:** Gunakan sparse atau linear attention untuk long sequences

### 6.2 Training Best Practices

1. **Learning Rate Warmup:** Mulai dengan learning rate kecil
2. **Label Smoothing:** Regularization technique untuk classification
3. **Dropout:** Apply pada attention dan feed-forward layers
4. **Layer Normalization:** Stabilize training
5. **Weight Initialization:** Xavier/He initialization

### 6.3 Debugging Strategies

1. Check attention weight sums (should be 1)
2. Verify mask is applied correctly
3. Monitor gradient norms
4. Visualize attention patterns
5. Start with small model and simple data

## 7 Evaluasi dan Penilaian

### 7.1 Komponen Penilaian

Komponen	Bobot	Keterangan
Kehadiran	10%	Partisipasi aktif
Percobaan 1-4	30%	Implementasi dan analisis
Latihan 1-3	20%	Problem solving
Latihan 4-6	25%	Advanced implementation
Laporan	15%	Dokumentasi dan analisis

Tabel 2: Distribusi penilaian praktikum

## 7.2 Kriteria Penilaian Laporan

### 1. Kelengkapan (25%)

- Semua percobaan dan latihan diselesaikan
- Code lengkap dan dapat dijalankan
- Visualisasi informatif

### 2. Analisis (35%)

- Pemahaman konsep mendalam
- Interpretasi hasil yang tepat
- Perbandingan metode yang komprehensif

### 3. Implementasi (25%)

- Code efficiency
- Proper error handling
- Code documentation

### 4. Presentasi (15%)

- Struktur laporan jelas
- Visualisasi berkualitas
- Penulisan profesional

## 8 Referensi

1. Vaswani, A., et al. (2017). "Attention Is All You Need." NeurIPS.
2. Bahdanau, D., Cho, K., & Bengio, Y. (2014). "Neural Machine Translation by Jointly Learning to Align and Translate." ICLR.
3. Devlin, J., et al. (2018). "BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding." NAACL.
4. Dosovitskiy, A., et al. (2021). "An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale." ICLR.
5. Brown, T., et al. (2020). "Language Models are Few-Shot Learners." NeurIPS.
6. Dive into Deep Learning. <https://d2l.ai>
7. The Illustrated Transformer. <http://jalammar.github.io/illustrated-transformer/>
8. Attention? Attention! <https://lilianweng.github.io/posts/2018-06-24-attention/>

## 9 Lampiran

### 9.1 Lampiran A: Cheat Sheet Formulas

#### 9.1.1 Attention Mechanisms

$$\text{Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V \quad (25)$$

$$\text{MultiHead}(Q, K, V) = \text{Concat}(h_1, \dots, h_n)W^O \quad (26)$$

$$h_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) \quad (27)$$

#### 9.1.2 Positional Encoding

$$PE_{(pos,2i)} = \sin \left( \frac{pos}{10000^{2i/d}} \right) \quad (28)$$

$$PE_{(pos,2i+1)} = \cos \left( \frac{pos}{10000^{2i/d}} \right) \quad (29)$$

#### 9.1.3 Layer Normalization

$$\text{LayerNorm}(x) = \gamma \odot \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} + \beta \quad (30)$$

## 9.2 Lampiran B: Common Hyperparameters

Parameter	Base	Large
$d_{model}$	512	1024
$d_{ff}$	2048	4096
$h$ (num heads)	8	16
$N$ (num layers)	6	12
Dropout	0.1	0.1

Tabel 3: Hyperparameter standar untuk Transformer

## 9.3 Lampiran C: Troubleshooting Guide

Problem	Solution
NaN loss	Check learning rate, gradient clipping, initialization
Poor convergence	Increase warmup steps, adjust LR schedule
Out of memory	Reduce batch size, use gradient accumulation
Attention weights not summing to 1	Check softmax axis, mask application
No learning	Check loss function, verify gradients flowing

Tabel 4: Common problems dan solutions